AI for PMs Starter Kit

By **Gonzalo Aguilar** · Free guide · No email required · Updated April 2026

AI can genuinely save a PM 3–5 hours a week on repetitive work. But most PMs either skip it entirely or try it once, get mediocre output, and conclude it doesn't work.

The secret is prompts. Vague inputs produce vague outputs. These 15 templates are the ones I actually use — built for real PM work, with enough specificity to get first-draft quality output you can actually use.

Everything here is copy-paste ready. Grab a prompt, fill in the brackets, and run it.

TABLE OF CONTENTS

1. 15 Prompt Templates for PM Work
2. Tool Comparison: ChatGPT vs. Claude vs. Gemini
3. Integrating AI into Your Sprint Cycle
4. Red Flags: When AI Output Needs Human Override

1. 15 Prompt Templates for PM Work

Templates use [BRACKETS] for the parts you fill in. Prompts are written for ChatGPT or Claude — both work well. Tips: always add context about your audience and product. More context = better output.

PROMPT 01 PRD / Spec Writing

PRD First Draft Generator

```
You are a senior product manager writing a PRD. Write a product requirements
document for the following feature:
```

```
Feature: [NAME OF FEATURE]
Product: [PRODUCT NAME — what it does, who uses it]
Problem: [THE USER PROBLEM THIS SOLVES IN 1-2 SENTENCES]
Target user: [WHO IS THIS FOR SPECIFICALLY]
Success metric: [HOW WILL YOU KNOW IF THIS WORKED]
Out of scope: [WHAT THIS FEATURE DOES NOT DO]

Write sections for: Problem Statement, Goals & Non-goals, User Stories (at
least 5), Functional Requirements, Edge Cases, Open Questions. Use clear
bullet points under each section. Keep it to 600-800 words — concise but
complete.
```

WHEN TO USE

Start of a new feature. Gives you a skeleton in ~2 minutes that would take 30+ minutes to draft from scratch.

IMPORTANT NOTE

Always review and add your own context, stakeholder inputs, and technical constraints. Never ship an AI PRD un-edited.

PROMPT 02

User Stories

User Story + Acceptance Criteria Generator

```
Write 5 user stories with acceptance criteria for the following feature:

Feature: [FEATURE NAME]
User persona: [WHO IS THE USER — e.g., "a freelancer who invoices clients
monthly"]
Core job to be done: [WHAT THEY'RE TRYING TO ACCOMPLISH]
Current pain: [WHAT'S FRUSTRATING ABOUT THE CURRENT EXPERIENCE]

Format each story as:
- As a [user], I want [action] so that [outcome]
- Acceptance Criteria: (3-5 specific, testable criteria using Given/When/Then
format)
- Edge case to consider: (1 edge case per story)

Make the acceptance criteria concrete and testable, not vague.
```

WHEN TO USE

PRO TIP

Sprint planning prep. Paste the output into Jira tickets as starting points.

Ask it to add a "negative story" — what the user should NOT be able to do. Prevents scope creep.

PROMPT 03

Competitive Analysis

Competitive Analysis Framework

```
Act as a product strategist. Analyze the competitive landscape for [YOUR
PRODUCT CATEGORY].

My product: [PRODUCT NAME + BRIEF DESCRIPTION]
My target customer: [WHO YOU SERVE]
Competitors to analyze: [LIST 3-5 COMPETITORS]

For each competitor, cover:
1. Core value proposition
2. How they position themselves (what they lead with)
3. Their apparent target customer
4. 3 things they do better than most
5. 3 observable weaknesses or gaps
6. Their pricing model (if known)

Then: Identify the top 2-3 whitespace opportunities — where are all these
competitors weak or silent in a way that matters to customers?

Note: flag any areas where your analysis may be incomplete or where I should
verify with primary research.
```

WHEN TO USE

Quarterly strategy reviews, new feature evaluation, investor prep. Use as a starting framework, then verify with actual product research.

IMPORTANT

AI competitive analysis is based on training data (often outdated). Always validate key claims by actually using the competitor products.

PROMPT 04

Data Analysis

Metric Interpretation + Next Steps

```
You are a product analyst. Help me interpret this data and suggest next
steps.

Metric I'm looking at: [METRIC NAME — e.g., "7-day retention rate"]
Current value: [NUMBER]
Previous period: [NUMBER]
Context: [WHAT'S HAPPENING IN THE PRODUCT — any recent releases, campaigns,
or changes]
Business context: [WHAT THIS METRIC MEANS FOR THE BUSINESS]

Please:
1. Identify the 3 most likely explanations for this change
2. For each explanation, tell me what data I should look at to confirm or
rule it out
3. Suggest 3 concrete next steps — one short-term investigation, one medium-
term test, one longer-term structural fix
4. Flag any data quality issues I should sanity-check before drawing
conclusions
```

WHEN TO USE

Weekly metric reviews, post-launch analysis, board prep. Helps you structure your thinking quickly.

PRO TIP

Paste in the actual data table for richer analysis. The more context you give, the better the output.

PROMPT 05

Stakeholder Communication

Difficult Stakeholder Email

```
Write a professional email to [STAKEHOLDER ROLE — e.g., "the VP of Sales"]
about [TOPIC].

Situation: [WHAT HAPPENED — be specific]
What I want them to know: [THE KEY MESSAGE]
What I want them to do: [CLEAR ASK, IF ANY]
Tone: [e.g., "clear and direct but not defensive" or "empathetic but firm"]
Things to avoid saying: [ANYTHING SENSITIVE OR OFF-LIMITS]

Keep it under 200 words. Use plain language — no corporate jargon. Lead with
the most important thing. End with a clear next step.
```

WHEN TO USE

Scope change announcements, timeline slips, pushing back on requests. Any time you need to write a careful email fast.

PRO TIP

Ask for 3 versions at different tone levels — then pick the one that fits the relationship.

PROMPT 06

User Research

User Interview Script Generator

```
Generate a 45-minute user interview script for the following research goal:

What I'm trying to learn: [RESEARCH QUESTION — e.g., "Why do users abandon
the checkout flow"]
User segment being interviewed: [WHO — e.g., "first-time buyers who haven't
converted"]
What I already know: [EXISTING HYPOTHESES OR DATA]
What I want to avoid: [LEADING QUESTIONS ABOUT, or topics to stay away from]

Include:
- 3 warm-up questions (to build rapport, not about the product)
- 8-10 core questions using open-ended, non-leading language
- 3 follow-up probes for each section ("tell me more about that", etc.)
- A closing question about what would make their experience better
- Timing estimate for each section

Flag any questions that might be leading and suggest alternatives.
```

WHEN TO USE

Discovery research, usability studies, post-launch learning. Saves 2 hours of script drafting.

IMPORTANT

Review for leading questions. AI will flag some but not always catch subtle bias. Get a colleague to review.

PROMPT 07

Strategy & Planning

OKR Writer

```
Write 3 OKRs for my team for the next quarter.

Team: [TEAM NAME / FUNCTION]
Company context: [WHAT THE COMPANY IS TRYING TO ACHIEVE THIS YEAR]
Team's primary responsibility: [WHAT YOUR TEAM OWNS]
Current biggest problem: [THE #1 THING HOLDING BACK PROGRESS]
North star metric: [THE METRIC THAT DEFINES SUCCESS FOR YOUR TEAM]

For each OKR:
- Objective: 1 sentence, qualitative, inspiring but grounded
- 3 Key Results: each must be specific, numeric, and achievable within the
quarter
- Confidence score (1-10): your gut check on whether this is achievable
- Risk: 1 main risk to hitting this KR

Avoid vanity metrics. Key Results should measure outcomes, not activity.
```

WHEN TO USE

Quarterly planning cycles. Get a draft in 5 minutes, then workshop it with the team.

PRO TIP

Ask it to critique your existing OKRs too — "which of these measure activity rather than outcomes?"

PROMPT 08

Communication

Release Notes for Multiple Audiences

```
Write release notes for the following product update in 3 formats:

Feature shipped: [WHAT SHIPPED]
Technical summary: [WHAT IT DOES IN TECHNICAL TERMS]
User benefit: [WHAT THE USER CAN NOW DO OR EXPERIENCE]
Who it affects: [WHICH USERS]

Write:
1. Internal Slack update for the team (casual, highlights what shipped and
any known issues, 50-75 words)
2. Customer-facing release note for your changelog (clear, benefit-focused,
no jargon, 75-100 words)
3. Executive summary bullet for the weekly leadership update (1-2 sentences,
metric-focused, ties to company goals)
```

WHEN TO USE

Every release. Saves 30-45 minutes of writing the same information in 3 different ways.

PRO TIP

Feed it your changelog entries from the last 3 releases to keep the tone consistent.

PROMPT 09

Prioritization

Backlog Prioritization Framework

```
Help me prioritize the following backlog items using the RICE framework.

Company stage: [EARLY STAGE / GROWTH / MATURE]
Current biggest company goal: [e.g., "Increase trial to paid conversion from
8% to 15%"]
Current biggest user complaint: [THE #1 THING USERS COMPLAIN ABOUT]

Backlog items to prioritize:
1. [ITEM 1 — brief description]
2. [ITEM 2]
3. [ITEM 3]
4. [ITEM 4]
5. [ITEM 5]

For each item, score it on:
- Reach: how many users/week would this impact?
- Impact: how much would this move the needle? (1=minimal, 3=medium,
5=massive)
- Confidence: how sure are we? (20%, 50%, 80%, 100%)
- Effort: weeks of engineering work (estimate)
- RICE Score: (Reach × Impact × Confidence) ÷ Effort

Rank them by RICE score. Flag any assumptions I should challenge.
```

WHEN TO USE

Sprint planning, roadmap reviews, stakeholder negotiation. Makes prioritization conversations data-driven.

IMPORTANT

The scores are only as good as your estimates. Use the output to structure discussion, not to make decisions mechanically.

PROMPT 10 Experimentation

A/B Test Design Brief

```
Design an A/B test for the following hypothesis.

Hypothesis: [YOUR HYPOTHESIS — e.g., "Showing social proof on the pricing
page will increase paid sign-ups"]
Current metric baseline: [CURRENT VALUE — e.g., "3.2% conversion from pricing
page to sign-up"]
Target metric: [WHAT YOU WANT TO MOVE]
Why we believe this will work: [THE REASONING BEHIND THE HYPOTHESIS]
What we'll change: [CONTROL VS. TREATMENT DESCRIPTION]
Risk of being wrong: [WHAT HAPPENS IF THE TEST DOESN'T WORK]

Design the test:
1. Recommended sample size and duration (calculate for 80% power, 5%
significance)
2. How to split traffic
3. What secondary metrics to watch (and what would indicate a tradeoff)
4. Pre-mortems: what could cause false positive or false negative results?
5. Decision criteria: what result do we need to call it a winner vs. loser
vs. inconclusive?
```

WHEN TO USE

Before every experiment. Catch flawed test designs before you run them for 6 weeks.

IMPORTANT

Verify the sample size math with a dedicated stats calculator. AI math can be off.

PROMPT 11 User Research

User Persona Builder

```
Create a detailed user persona based on the following research inputs.

Product: [PRODUCT NAME AND PURPOSE]
Research inputs: [PASTE IN: key quotes from interviews, survey data,
behavioral patterns, support tickets — whatever you have]
Segment focus: [WHICH USER SEGMENT — e.g., "power users who use the product
daily"]
```

```
Build a persona that includes:
- Name, age, role, and context (make it realistic, not a stereotype)
- Primary goal when using this product
- The job to be done (what they're really trying to accomplish)
- Top 3 frustrations with the current experience
- Top 3 things they love or that work well
- Quote that captures their mindset (based on research, not invented)
- Decision-making style: how do they evaluate and choose tools?
- What success looks like for them

Flag which elements are well-supported by the data vs. which are inferred.
```

WHEN TO USE

After user research rounds, onboarding new team members, refocusing a product on a specific segment.

PRO TIP

Ground it in real quotes. AI personas without research input are fictional and can mislead.

PROMPT 12

Stakeholder Communication

Executive Briefing / Status Update

```
Write a crisp executive status update for the following situation.

Audience: [WHO IS READING — e.g., "VP of Product + CFO in the weekly
leadership sync"]
Project/Initiative: [WHAT YOU'RE UPDATING THEM ON]
Current status: [GREEN / YELLOW / RED and why]
Key accomplishment since last update: [WHAT SHIPPED OR GOT DONE]
Current blocker or risk: [IF ANY — be specific, not vague]
What you need from them: [DECISION, APPROVAL, OR AWARENESS — be explicit]
Relevant metrics: [ANY NUMBERS THAT SHOW PROGRESS OR LACK THEREOF]

Write in this format:
- Status: [color] — [one sentence explanation]
- This week: [2-3 bullet points, past tense, outcome-focused]
- Risk: [1 sentence if applicable, with mitigation]
- Ask: [clear single ask, or "no action needed"]

Under 150 words total. No fluff.
```

WHEN TO USE

Weekly syncs, board decks, all-hands updates. Executives read fast — this format respects that.

PRO TIP

Always include the "ask" even if it's "no action needed." Removes ambiguity.

PROMPT 13 Discovery

Jobs-to-Be-Done Analysis

```
Apply the Jobs-to-Be-Done framework to help me understand why users are
"hiring" my product.

Product: [PRODUCT NAME AND BRIEF DESCRIPTION]
Known use cases: [HOW USERS DESCRIBE USING IT]
Data I have: [ANY USAGE PATTERNS, COHORTS, OR SEGMENTS I CAN SEE]
Problem I'm trying to solve: [WHAT STRATEGIC QUESTION THIS ANALYSIS IS FOR]

Using JTBD methodology, analyze:
1. Functional job: what practical task are users trying to complete?
2. Emotional job: how do they want to feel (or avoid feeling) when they use
this?
3. Social job: how do they want to be perceived by others when using this?
4. For each of the 3 job layers: what are the underserved outcomes — where
are current solutions (including mine) falling short?
5. Which competitor or substitute solution is being "fired" when a user hires
my product?
6. Implications: what features or positioning changes does this analysis
suggest?
```

WHEN TO USE

Strategic reviews, positioning work, discovery sprints, new market entry analysis.

PRO TIP

Feed in real customer quotes for the best output. The emotional and social layers especially benefit from real language.

PROMPT 14 Retrospectives

Post-Mortem Facilitator

```
Help me structure a blameless post-mortem for the following incident or
failed project.

What happened: [DESCRIBE THE INCIDENT OR FAILURE — timeline, impact, who was
affected]
What we expected to happen: [THE INTENDED OUTCOME]
Timeline of events: [KEY MOMENTS — what decisions were made and when]
Immediate causes: [WHAT DIRECTLY CAUSED THE PROBLEM]
Contributing factors: [SYSTEM, PROCESS, OR COMMUNICATION ISSUES THAT ALLOWED
IT]

Facilitate the post-mortem by:
1. Reframing each cause in blameless terms (system/process failure, not
personal failure)
2. Asking "5 Whys" for the top 2-3 contributing factors
3. Generating action items in this format: [What], [Who owns it], [By when],
[How we'll know it worked]
4. Identifying 1-2 early warning indicators we could monitor to catch this
earlier next time
5. Drafting a 3-sentence summary we can share with stakeholders
```

WHEN TO USE

After any significant incident, missed deadline, or failed launch. Gets past blame-finding to actual fixes.

IMPORTANT

Still run it live with the team — AI gives you structure, but the conversation is where learning happens.

PROMPT 15

Strategy & Planning

Outcome-Based Roadmap Draft

```
Help me draft a 6-month outcome-based roadmap.

Product: [PRODUCT NAME AND BRIEF DESCRIPTION]
Company goal this year: [THE #1 BUSINESS OBJECTIVE]
Current product state: [WHERE THE PRODUCT IS TODAY — key capabilities and
gaps]
Known user problems (in priority order): [LIST THE TOP 3-5 PROBLEMS YOU KNOW
ABOUT]
Technical debt or platform work to account for: [ANY MUST-DO INFRASTRUCTURE
WORK]
```

```
Team size: [NUMBER OF ENGINEERS, DESIGNERS, AND PMs]

Draft a roadmap organized by outcome (not feature):
- Month 1-2: Foundation (what must be true for everything else to work)
- Month 3-4: Acceleration (what drives the core metric)
- Month 5-6: Expansion (what opens new opportunity)

For each phase:
- The outcome we're trying to achieve (measurable)
- 3-5 bets that could achieve that outcome
- Dependencies and risks
- What "done" looks like

Format it so I can drop it into a slide or doc immediately.
```

WHEN TO USE

Planning cycles, new PM onboarding, alignment conversations with leadership.

PRO TIP

Run this, then ask it to "steelman the 3 biggest criticisms of this roadmap." The pushback is useful.

2. Tool Comparison: ChatGPT vs. Claude vs. Gemini for PM Work

Task	ChatGPT (GPT-4o)	Claude (Sonnet/Opus)	Gemini (Advanced)
PRD & spec writing	Excellent. Strong structure, good at following format templates	Excellent. Especially good at nuance and avoiding generic language	Good. Can be more verbose; needs tighter prompts
User stories + ACs	Very good. Follows Given/When/Then well	Very good. Strong at edge cases when prompted	Good. Sometimes misses testability requirement

Task	ChatGPT (GPT-4o)	Claude (Sonnet/Opus)	Gemini (Advanced)
Data analysis & interpretation	Good. GPT-4o can run code and analyze actual CSV data	Good at frameworks; weaker at running live calculations	Good integration with Google Sheets; strong for Google-ecosystem users
Long document analysis	Good; 128K context window	Best-in-class; 200K context. Excellent at following instructions across long docs	Very good; 1M context window (Gemini 1.5)
Stakeholder emails	Very good. Versatile tone control	Excellent. Tone nuance is a Claude strength	Good but can sound slightly stiff
Competitive research	Good but training cutoff applies	Same limitation but strong structured analysis	Can access live Google search data (with web grounding)
Code-adjacent tasks (tech specs)	Best. GPT-4o is strong at technical writing and code	Very good at technical writing; excellent at code	Good; especially for Google Cloud / GCP context
Price (as of mid-2025)	$20/mo consumer; API pricing varies	$20/mo consumer; API pricing varies	$20/mo (Google One AI Premium)

> **MY RECOMMENDATION**
>
> **For pure PM writing work** (PRDs, emails, user stories): Claude is slightly better at tone and nuance.
> **For data analysis**: ChatGPT with code interpreter, or Gemini if you're Google-native.
> **For long document analysis**: Claude (200K context) or Gemini 1.5 (1M context).
> **For current information**: Gemini (web grounding) or ChatGPT with browsing enabled.

---

3. Integrating AI into Your Sprint Cycle

The goal isn't to use AI on every task — it's to identify where you have repetitive high-effort work and let AI take the first draft. Here's how it fits across a typical sprint:

Before Sprint Planning

- Use Prompt 09 to draft RICE scores for backlog items — paste into planning doc for discussion
- Use Prompt 07 to draft OKRs for the quarter — workshop them with the team
- Use Prompt 03 to refresh competitive context if new features are being planned

During the Sprint

- Use Prompt 01 (PRD) and Prompt 02 (User Stories) for any new feature work — start with AI draft, edit, share
- Use Prompt 05 for any difficult stakeholder communication — never send un-edited
- Use Prompt 04 if you're investigating a metric movement mid-sprint

Sprint Review / Retrospective

- Use Prompt 08 to generate release notes across 3 audiences in one pass
- Use Prompt 12 for executive status updates
- Use Prompt 14 for post-mortems if something went wrong

Discovery and Research Work

- Use Prompt 06 (interview script) before any user research session
- Use Prompt 11 to synthesize research into personas
- Use Prompt 13 for JTBD analysis when doing strategic reviews

> **THE RULE I FOLLOW**
>
> AI handles the first draft. I handle the judgment. I've never shipped a PRD, email, or roadmap that came directly from AI without human review. The output gets me to 70% faster — the final 30% is still mine.

4. Red Flags: When AI Output Needs Human Override

⚠ **The output sounds confident but cites specific numbers you haven't verified.** Always check any statistics, market size claims, or competitor data. AI is trained on text that contains errors.

⚠ **The tone is too smooth and generic.** Good communication requires voice and specificity. If it reads like a template rather than a message from a real person, edit it heavily before sending.

⚠ **You're using AI for a decision, not a document.** AI is excellent at producing documents. It's not a substitute for judgment on calls like "should we build this feature" or "is this user research finding significant?"

⚠ **Legal, compliance, or regulated-industry content.** AI doesn't know the specific legal or regulatory environment you're operating in. Get a human expert involved before anything touches legal or compliance territory.

⚠ **The output doesn't match what users actually said.** If an AI-generated persona or problem statement contradicts your qualitative research, trust the research. AI summarizes text patterns — you have direct access to your users.

⚠ **It's writing about interpersonal situations.** Conflict resolution, performance conversations, difficult feedback — these require human judgment about relationship dynamics. Use AI to prepare, not to replace the conversation.

⚠ **The output is technically correct but misses the real problem.** AI solves the problem you described. If you described the wrong problem, the output will be well-executed but irrelevant. Your job is to make sure the prompt captures the real situation.

Gonzalo Aguilar is a Senior Product Manager specializing in growth, identity, and onboarding in regulated markets. He writes about product strategy, conversion optimization, and practical AI workflows at gonzalo.tech.